

Lecture 2:

Optimal Classifier

Spring 2026


10 ECTS Credits

Department of Information and Communication Technology

Faculty of Engineering and Science


Recap

- Pattern recognition pipeline
- Single feature classification
- Decision threshold
- Two feature classification
- Some problems in finding an optimal model
 - Outliers and missing values
 - Feature extraction
 - Costs and risks
- Learning
 - Underfitting vs. Overfitting
 - Types of Learning



The problem

- Problem $\rightarrow$ feature vector and target
- X, Y relationship is stochastic in nature
- The relationship can only be specified by a joint feature-target distribution, $P_{X,Y}$.
- However:
 - Y is dependent on hidden/latent factors.
 - There are measurement noise in X .



Prediction

- Supervised learning objective: **predict** Y given X .
- Predictor, $\psi: \mathbb{R}^d \rightarrow \mathbb{R}$, such that $\psi(X)$ predicts Y and is designed using either:
 - Direct knowledge about $P_{X,Y}$, which can yield $\psi^*(X)$.
 - Indirect knowledge about $P_{X,Y}$ through an independent and identically distributed (i.i.d.) sample $S_n = \{(X_1, Y_1), \dots, (X_n, Y_n)\}$
- A **loss function** $\ell: \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$ defines the validity or optimality of a predictor.
- Loss functions like **quadratic loss**, **absolute difference loss**, **misclassification loss**.
- Since we deal with stochastic model, **prediction error** is given by:

$$L[\psi] = E[\ell(Y, \psi(X))]$$

- The optimal predictor:

$$\psi^* = \arg \min_{\psi \in \mathcal{P}} L[\psi]$$

Recall: Expectation (Expected Value)

- Is expectation operation same as the average/mean?
- Expectation is the theoretical average of a random variable across its entire probability distribution:

$$E[X] = \begin{cases} \sum_x x \cdot P(X = x), & \text{(discrete)} \\ \int_{-\infty}^{\infty} x \cdot f(x) dx, & \text{(continuous)} \end{cases}$$

- $P(X = x)$ is the probability of $X = x$.
- $f(x)$ is the probability density function (PDF) of the continuous random variable.



Recall: Expectation (Expected Value)

- $X \rightarrow$ Outcome of rolling a fair die

$$E[X] = \sum_x x \cdot P(X = x)$$



$$E[X] = 1 \cdot \frac{1}{6} + 2 \cdot \frac{1}{6} + 3 \cdot \frac{1}{6} + 4 \cdot \frac{1}{6} + 5 \cdot \frac{1}{6} + 6 \cdot \frac{1}{6} = 3.5$$

- If you roll $\{5, 2, 6, 2, 2, 1, 2, 3, 6, 1\}$, sample mean (average) over $N = 10$ rolls:
 $\bar{x} = (5 + 2 + 6 + 2 + 2 + 1 + 2 + 3 + 6 + 1)/10 = 3.0$
- We can say the sample mean is 3.0 with a distance of 0.5 from the expected value of 3.5.
- However, as $N \rightarrow \infty$, $\bar{x} \rightarrow E[X]$.



Complexity Trade-offs

- **Curse of dimensionality** or Hughes Phenomenon:
 - Accuracy increases and then decreases as number of features increases.
 - The peaks shift right with increased sample size
- Seen in **scissors plot**:
 - Complex classifier: error converges to an optimal error with sample size.
 - Simple classifier: performs better under small sample sizes.
 - A problem dependent N_0 , under which the simpler classifier performs better.



Classification without features

- Joint feature-label distribution $P_{X,Y}$ available → **optimal classifier** can be obtained.
- Simple case: a binary classifier with no features, $\hat{Y}$, such that $Y \in \{0,1\}$
- $\hat{Y}$ assign the label of the most common class.
- Class **prior probabilities/prevalences**: $P(Y = 0)$ and $P(Y = 1)$.
- Assigning a label closest to the mean.
- Rare disease example

$$\hat{Y} = \begin{cases} 1, P(Y = 1) > P(Y = 0), \\ 0, P(Y = 1) \leq P(Y = 0). \end{cases}$$

$$P(Y = 0) + P(Y = 1) = 1$$

$$\hat{Y} = 1, \text{ iff } P(Y = 1) > 0.5$$

$$E[Y] = P(Y = 1)$$

$$\varepsilon[\hat{Y}] = P(\hat{Y} \neq Y) = \min\{P(Y = 1), P(Y = 0)\}$$

Recall: Probability Density Functions (PDFs)

- Let us consider a task with a single continuous feature $x \in \mathbb{R}$ and plot its PDF, $p(x)$.


- What does this plot show?

- What is $P(x = 2)$?

$$P(x = 2) = 0$$


- What is $P(4 \leq x \leq 6)$?

$$P(4 \leq x \leq 6) = \int_E p(x)dx = 0.321$$



Recall: Probability Density Functions (PDFs)

- What are **class-conditional densities**?
- If I already know this instance belongs to class 0 (or 1), what feature values am I likely to see?
 - $p(x | Y = 0)$ → Distribution of the feature among class 0 instances.
 - $p(x | Y = 1)$ → Distribution of the feature among class 1 instances.
- However, class-conditional densities doesn't account for how often each class occurs.



Recall: Probability Density Functions (PDFs)

- The class-conditional densities are weighted by the priors.

$$p(x, Y = 0) = P(Y = 0)p(x | Y = 0)$$

$$p(x, Y = 1) = P(Y = 1)p(x | Y = 1)$$


- These curves represent the actual contribution of classes to the overall feature distribution.

- If we add these together, we get:

$$P(Y = 0)p(x | Y = 0) + P(Y = 1)p(x | Y = 1) = p(x)$$


- This is known as the **marginal PDF** of the feature.

$$p(x) = \sum_{Y \in \{0,1\}} p(x, y)$$



Classification with features

- Each $\mathbf{X} \in \mathbb{R}^d$ is a **feature vector**,
 - where d , is the **dimension**,
 - and $\mathbb{R}^d$ is called the **feature space**.
- We assume all features are real-valued and continuous.
- **Class conditional densities** $p(\mathbf{x}|Y = 0)$ and $p(\mathbf{x}|Y = 1)$.
- When $E = \mathbb{R}^d$, LHS becomes prior probabilities.
- Marginal PDF $p(\mathbf{x})$:
- $P_{\mathbf{X},Y}$ is fully specified by the class-conditional densities weighted by prior probabilities.



$$\mathbf{X} = [X_1, X_2, \dots, X_d] \in \mathbb{R}^d$$

$$\mathcal{C} = \{0,1\}$$

$$P(\mathbf{X} \in E, Y = c) = P(Y = c) \int_E p(\mathbf{x}|Y = c) d\mathbf{x}$$
$$E \subseteq \mathbb{R}^d$$

$$p(\mathbf{x}) = P(Y = 0)p(\mathbf{x}|Y = 0) + P(Y = 1)p(\mathbf{x}|Y = 1)$$

Cookie Quality Control

- In a bakery, 5% cookies are perfectly baked (95% are not).
- A human inspector flags
 - 85% of perfectly baked cookies correctly,
 - 12% of imperfect cookies wrong.
- The question:
If the inspector says a cookie is perfect, what's the actual chance it really is perfect?
- Answer: approx. **27%**. Why?



Cookie Quality Control

	Actually Perfect (5%)	Not Perfect (95%)
Inspector Says "Perfect"	85% (True Positive)	12% (False Positive)
Inspector Says "Not Perfect"	15% (False Negative)	88% (True Negative)


If the inspector says a cookie is perfect, what's the actual chance it really is perfect?

- Let us say, 10000 cookies are tested a month:

$$P(\text{Perfect} \mid \text{Says Perfect}) = \frac{425}{425 + 1140} \approx 27.16\%$$

$$= \frac{P(\text{Perfect} \mid \text{Says Perfect})}{P(\text{Says Perfect})} = \frac{P(\text{Says Perfect} \mid \text{Perfect}) \times P(\text{Perfect})}{P(\text{Says Perfect})}$$

- Why the very small value?



Bayes' Theorem

- The theorem states that:

$$P(H | E) = \frac{P(E | H)P(H)}{P(E)} = \frac{P(E | H)P(H)}{P(E | H)P(H) + P(E | \text{not } H)P(\text{not } H)}$$

$H \rightarrow$ Hypothesis, $E \rightarrow$ Event

- Or more generally for two events A and B and $P(B) \neq 0$:

$$P(A | B) = \frac{P(B | A)P(A)}{P(B)}$$

- $P(A | B) \rightarrow$ Conditional probability of A occurring given B is true.
- $P(B | A) \rightarrow$ Conditional probability of B occurring given A is true.
- $P(A)$ and $P(B)$ are probabilities of A and B occurring independently.


Posterior probabilities

- Assuming $p(\mathbf{x}) > 0$:
- **Posterior probability** of Y given $X = \mathbf{x}$:

$$P(Y = 0|X = \mathbf{x}) = \frac{P(Y = 0)p(\mathbf{x}|Y = 0)}{p(\mathbf{x})} = \frac{P(Y = 0)p(\mathbf{x}|Y = 0)}{P(Y = 0)p(\mathbf{x}|Y = 0) + P(Y = 1)p(\mathbf{x}|Y = 1)}$$

$$P(Y = 1|X = \mathbf{x}) = \frac{P(Y = 1)p(\mathbf{x}|Y = 1)}{p(\mathbf{x})} = \frac{P(Y = 1)p(\mathbf{x}|Y = 1)}{P(Y = 0)p(\mathbf{x}|Y = 0) + P(Y = 1)p(\mathbf{x}|Y = 1)}$$

- Value ranges between 0 and 1, and
- Sum up to 1, for the entire feature space.
- Thus, for a binary classification, we can use the **posterior probability function**:
$$\eta(\mathbf{x}) = E[Y|X = \mathbf{x}] = P(Y = 1|X = \mathbf{x}), \mathbf{x} \in \mathbb{R}^d$$




Classifiers and Classification Error

- A **classifier** is a function $\psi: \mathbb{R}^d \rightarrow \{0,1\}$ that divides the feature space $\mathbb{R}^d$ into two sets:
 - a 0-decision region, and
 - a 1-decision regionby a **decision boundary**, x_0 .
- Classification error is the probability of misclassification:

$$\varepsilon[\psi] = P(\psi(X) \neq Y)$$

- Main criterion for performance in classification, determined by $P_{X,Y}$



Class-specific Errors

- **Class specific error rates** of ψ measures misclassification rates within each class:

- **False Positive Rate**

$$\varepsilon^0[\psi] = P(\psi(X) = 1|Y = 0) = \int_{\{x|\psi(x)=1\}} p(x|Y=0)dx$$

- **False Negative Rate**

$$\varepsilon^1[\psi] = P(\psi(X) = 0|Y = 1) = \int_{\{x|\psi(x)=0\}} p(x|Y=1)dx$$

- **Total error (1):**


$$\varepsilon[\psi] = P(\psi(X) \neq Y) = P(\psi(X) = 1, Y = 0) + P(\psi(X) = 0, Y = 1)$$

$$= P(Y = 0)P(\psi(X) = 1|Y = 0) + P(Y = 1)P(\psi(X) = 0|Y = 1)$$

$$= P(Y = 0) \int_{\{x|\psi(x)=1\}} p(x|Y=0)dx + P(Y = 1) \int_{\{x|\psi(x)=0\}} p(x|Y=1)dx$$

$$= (1 - n)\varepsilon^0[\psi] + n\varepsilon^1[\psi]$$

- A linear combination of class-specific error rates weighted by class priors.



Conditional Classification Error

- **Conditional error** of a classifier:

$$\varepsilon[\psi | \mathbf{X} = \mathbf{x}] = P(\psi(\mathbf{X}) \neq Y | \mathbf{X} = \mathbf{x})$$

- Using the **Law of Total Probability for random variables**, classification error is the expected conditional classification error over the feature space (2):


$$\varepsilon[\psi] = E[\varepsilon(\psi | \mathbf{X} = \mathbf{x})] = \int_{\mathbf{x} \in \mathbb{R}^d} \varepsilon[\psi | \mathbf{X} = \mathbf{x}] p(\mathbf{x}) d\mathbf{x}$$

- The conditional classification error $\varepsilon[\psi | \mathbf{X} = \mathbf{x}]$ is determined by the posterior probability function $\eta(\mathbf{x})$ as (3):

$$\begin{aligned} \varepsilon[\psi | \mathbf{X} = \mathbf{x}] &= P(\psi(\mathbf{X}) = 0, Y = 1 | \mathbf{X} = \mathbf{x}) + P(\psi(\mathbf{X}) = 1, Y = 0 | \mathbf{X} = \mathbf{x}) \\ &= I_{\psi(\mathbf{x})=0} P(Y = 1 | \mathbf{X} = \mathbf{x}) + I_{\psi(\mathbf{x})=1} P(Y = 0 | \mathbf{X} = \mathbf{x}) \\ &= I_{\psi(\mathbf{x})=0} \eta(\mathbf{x}) + I_{\psi(\mathbf{x})=1} (1 - \eta(\mathbf{x})) = \begin{cases} \eta(\mathbf{x}), & \text{if } \psi(\mathbf{x}) = 0 \\ 1 - \eta(\mathbf{x}), & \text{if } \psi(\mathbf{x}) = 1 \end{cases} \end{aligned}$$

- Recall, for the above derivation:

- $\eta(\mathbf{x}) = P(Y = 1 | \mathbf{X} = \mathbf{x})$: Probability that the true label is 1, given $\mathbf{X} = \mathbf{x}$.
- $1 - \eta(\mathbf{x}) = P(Y = 0 | \mathbf{X} = \mathbf{x})$: Probability that the true label is 0, given $\mathbf{X} = \mathbf{x}$.



Bayes Classifier

- A **Bayes classifier** minimizes the classification error:

$$\psi^* = \arg \min_{\psi \in \mathcal{C}} P(\psi(\mathbf{X}) \neq Y)$$

- There maybe an infinite number of Bayes classifier for a given problem.
- **Theorem (4):** The classifier

$$\psi^*(\mathbf{x}) = \arg \max_i P(Y = i \mid \mathbf{X} = \mathbf{x}) = \begin{cases} 1, & \eta(\mathbf{x}) > 0.5 \\ 0, & \text{otherwise} \end{cases}$$

for $\mathbf{x} \in \mathbb{R}^d$ is a Bayes Classifier.

- Only the knowledge of the posterior probability function is necessary to obtain a Bayes Classifier.

Theorem: Proof

- For all $\psi \in \mathcal{C}$, we want to show that

$$\varepsilon[\psi] \geq \varepsilon^*[\psi]$$

- From (2)

$$\varepsilon[\psi \mid \mathbf{X} = \mathbf{x}] \geq \varepsilon[\psi^* \mid \mathbf{X} = \mathbf{x}], \forall \mathbf{x} \in \mathbb{R}^d$$

- Or,

$$\varepsilon[\psi \mid \mathbf{X} = \mathbf{x}] - \varepsilon[\psi^* \mid \mathbf{X} = \mathbf{x}] \geq 0, \forall \mathbf{x} \in \mathbb{R}^d$$

- Using (3),

$$\begin{aligned} & \varepsilon[\psi \mid \mathbf{X} = \mathbf{x}] - \varepsilon[\psi^* \mid \mathbf{X} = \mathbf{x}] \\ &= I_{\psi(\mathbf{x})=0}\eta(\mathbf{x}) + I_{\psi(\mathbf{x})=1}(1 - \eta(\mathbf{x})) - \left(I_{\psi^*(\mathbf{x})=0}\eta(\mathbf{x}) + I_{\psi^*(\mathbf{x})=1}(1 - \eta(\mathbf{x})) \right) \\ &= \eta(\mathbf{x})(I_{\psi(\mathbf{x})=0} - I_{\psi^*(\mathbf{x})=0}) + (1 - \eta(\mathbf{x}))(I_{\psi(\mathbf{x})=1} - I_{\psi^*(\mathbf{x})=1}) \end{aligned}$$

Theorem: Proof

- Let us look at the possibilities in making up these differences:

$$I_{\psi(x)=0} - I_{\psi^*(x)=0}, \quad I_{\psi(x)=1} - I_{\psi^*(x)=1}$$

- Exhausting all the possibilities for $\psi(x)$ and $\psi^*(x)$:

$$I_{\psi(x)=0} - I_{\psi^*(x)=0} = -(I_{\psi(x)=1} - I_{\psi^*(x)=1})$$


$\psi(x)$	$\psi^*(x)$	$I_{\psi(x)=0}$	$I_{\psi^*(x)=0}$	$I_{\psi(x)=0} - I_{\psi^*(x)=0}$	$I_{\psi(x)=1}$	$I_{\psi^*(x)=1}$	$I_{\psi(x)=1} - I_{\psi^*(x)=1}$
0	0	1	1	0	0	0	0
0	1	1	0	+1	0	1	-1
1	0	0	1	-1	1	0	+1
1	1	0	0	0	1	1	0

Theorem: Proof

- Substituting $I_{\psi(x)=0} - I_{\psi^*(x)=0} = -(I_{\psi(x)=1} - I_{\psi^*(x)=1})$,
$$\begin{aligned}\varepsilon[\psi \mid \mathbf{X} = \mathbf{x}] - \varepsilon[\psi^* \mid \mathbf{X} = \mathbf{x}] &= \eta(\mathbf{x})(I_{\psi(x)=0} - I_{\psi^*(x)=0}) - (1 - \eta(\mathbf{x}))(I_{\psi(x)=0} - I_{\psi^*(x)=0}) \\ &= (2\eta(\mathbf{x}) - 1)(I_{\psi(x)=0} - I_{\psi^*(x)=0})\end{aligned}$$
- There are only two possibilities:
 - $\eta(\mathbf{x}) > \frac{1}{2} \rightarrow$ Both terms in the parentheses are nonnegative.
 - $\eta(\mathbf{x}) \leq \frac{1}{2} \rightarrow$ Both terms in the parentheses are nonpositive.
- Thus, in both cases the product is nonnegative meaning:
$$\varepsilon[\psi \mid \mathbf{X} = \mathbf{x}] - \varepsilon[\psi^* \mid \mathbf{X} = \mathbf{x}] \geq 0, \forall \mathbf{x} \in \mathbb{R}^d$$

Bayes Classifier (Contd.)

- Optimal decision boundary is the set:
$$\{\mathbf{x} \in R^d \mid \eta(\mathbf{x}) = 0.5\}$$
- Bayes classifier is defined pointwise to minimize the error $\varepsilon[\psi \mid \mathbf{X} = \mathbf{x}]$ at each value of $\mathbf{x} \in R^d$, which minimizes the error globally.
- Bayes classifier is both local and global optimal predictor.



Bayes classifier (Contd.)

- Bayes classifier theorem for $x \in R^d$ can be restated as:

$$\psi^*(x) = \begin{cases} 1, & \eta(x) > 1 - \eta(x) \\ 0, & \text{otherwise} \end{cases}$$

- Considering posterior probability equation:

$$\psi^*(x) = \begin{cases} 1, & P(Y = 1)p(x | Y = 1) > P(Y = 0)p(x | Y = 0) \\ 0, & \text{otherwise} \end{cases}$$

- Bayes classifier can be determined by comparing at each point in feature space the weighted class-conditional densities.
- The above equation can be rewritten as:

$$\psi^*(x) = \begin{cases} 1, & D^*(x) > k^* \\ 0, & \text{otherwise} \end{cases}$$

where the **optimal discriminant** $D^*: R^d \rightarrow R$ is given by

$$D^*(x) = \ln \frac{p(x | Y = 1)}{p(x | Y = 0)}$$

for $x \in R^d$, with **optimal threshold** k^* given by

$$k^* = \ln \frac{P(Y = 0)}{P(Y = 1)}$$

Bayes Classifier

- A **Bayes classifier** minimizes the classification error:
$$\psi^* = \arg \min_{\psi \in \mathcal{C}} P(\psi(X) \neq Y)$$
- There may be an infinite number of Bayes classifier for a given problem.
- Theorem (4):** The classifier
$$\psi^*(x) = \arg \max_i P(Y = i | X = x) = \begin{cases} 1, & \eta(x) > 0.5 \\ 0, & \text{otherwise} \end{cases}$$
for $x \in \mathbb{R}^d$ is a Bayes Classifier.
- Only the knowledge of the posterior probability function is necessary to obtain a Bayes Classifier.



Lecture 2: Optimal Classifier


20

Bayes classifier (Contd.)

- Bayes Classifier:

$$\psi^*(\mathbf{x}) = \begin{cases} 1, & D^*(\mathbf{x}) > k^* \\ 0, & \text{otherwise} \end{cases}, \quad D^*(\mathbf{x}) = \ln \frac{p(\mathbf{x} | Y = 1)}{p(\mathbf{x} | Y = 0)}, \quad k^* = \ln \frac{P(Y = 0)}{P(Y = 1)}$$

- $D^*(\mathbf{x})$ is called the **log-likelihood ratio** or **optimal discriminant function**:
 - Measures how much more likely $\mathbf{x}$ is under class 1 vs. class 0.
 - With $D^*(\mathbf{x}) > k^*$, $\mathbf{x}$ provides enough evidence to predict class 1.
- k^* adjusts for **class imbalance**:
 - If Class 0 is more common ($P(Y = 0) > P(Y = 1)$), $k^* > 0$. You need a stronger evidence ($D^*(\mathbf{x})$) to predict class 1.
 - If $P(Y = 0) = P(Y = 1)$, $k^* = 0$, we compare likelihoods directly.



Bayes Error

- Given a Bayes classifier ψ^* , the error $\varepsilon^* = \varepsilon[\psi^*]$ is known as **Bayes error**.
- All Bayes classifiers share the same Bayes error: the lower bound on the classification error that can be achieved in a given problem.
- Should be sufficiently small to design a good classifier from data.
- **Theorem:** In two-class classification problem, the Bayes error is given by:
$$\varepsilon^* = E[\min\{\eta(\mathbf{X}), 1 - \eta(\mathbf{X})\}]$$

Further, the maximum value Bayes error can take is 0.5.

Bayes Error – Proof

- From (2), (3), and (4):

$$\begin{aligned}\varepsilon^* &= E[\varepsilon(\psi^* | \mathbf{X} = \mathbf{x})] = E[I_{\psi^*(\mathbf{x})=0}\eta(\mathbf{x}) + I_{\psi^*(\mathbf{x})=1}(1 - \eta(\mathbf{x}))] \\ &= E[I_{\eta(\mathbf{x}) \leq 1 - \eta(\mathbf{x})}\eta(\mathbf{X}) + I_{\eta(\mathbf{x}) > 1 - \eta(\mathbf{x})}(1 - \eta(\mathbf{X}))] = E[\min\{\eta(\mathbf{X}), 1 - \eta(\mathbf{X})\}]\end{aligned}$$

- Applying the identity $\min\{a, 1 - a\} = \frac{1}{2} - \frac{1}{2}|2a - 1|, \forall 0 \leq a \leq 1$

$$\varepsilon^* = \frac{1}{2} - \frac{1}{2}E[|2\eta(\mathbf{X}) - 1|]$$

$$\Rightarrow \varepsilon^* \leq \frac{1}{2}$$

- The maximum optimal error in a two-class classification problem is 0.5.
- Also,

$$\varepsilon^* = \frac{1}{2} \Leftrightarrow E[|2\eta(\mathbf{X}) - 1|] = 0 \Leftrightarrow \eta(\mathbf{X}) = \frac{1}{2} \text{ with probability 1.}$$

- A problem, where Bayes error is 0.5 or very close to it, is hopeless.

Bayes Error – Proof (Contd.)


- Recall that for $x \in R^d$:

$$\psi^*(\mathbf{x}) = \begin{cases} 1, & P(Y = 1)p(\mathbf{x} | Y = 1) > P(Y = 0)p(\mathbf{x} | Y = 0) \\ 0, & \text{otherwise} \end{cases}$$

- Applying (1) above:

$$\begin{aligned} \varepsilon[\psi^*] &= P(Y = 0)\varepsilon^0[\psi^*] + P(Y = 1)\varepsilon^1[\psi^*] \\ &= \int_{\{P(Y=1)p(\mathbf{x}|Y=1) > P(Y=0)p(\mathbf{x}|Y=0)\}} P(Y = 0)p(\mathbf{x}|Y = 0)d\mathbf{x} \\ &\quad + \int_{\{P(Y=1)p(\mathbf{x}|Y=1) \leq P(Y=0)p(\mathbf{x}|Y=0)\}} P(Y = 1)p(\mathbf{x}|Y = 1)d\mathbf{x} \end{aligned}$$

- Bayes error is the sum of integrals of prior-weighted densities over the regions of overlap between them.



Bayes Error – Transformed feature vector

- Feature transformation: preprocessing, normalization, dimensionality reduction, etc.
- **Theorem:** Given a transformation between feature spaces $t: \mathbb{R}^p \rightarrow \mathbb{R}^d$ such that it transforms the original feature vector $\mathbf{X} \in \mathbb{R}^p$ as $\mathbf{X}' = t(\mathbf{X}) \in \mathbb{R}^d$ and the Bayes errors $\varepsilon^*(\mathbf{X}, Y)$ and $\varepsilon^*(\mathbf{X}', Y)$:

$$\varepsilon^*(\mathbf{X}', Y) \geq \varepsilon^*(\mathbf{X}, Y)$$

with equality if t is reversible.

Minimax classification

- Bayes classifier:

$$\psi^*(\mathbf{x}) = \begin{cases} 1, & D^*(\mathbf{x}) > k^* \\ 0, & \text{otherwise} \end{cases}$$

where the **optimal discriminant** $D^*: \mathbb{R}^d \rightarrow \mathbb{R}$ is given by

$$D^*(\mathbf{x}) = \ln \frac{p(\mathbf{x} | Y = 1)}{p(\mathbf{x} | Y = 0)}$$

for $\mathbf{x} \in \mathbb{R}^d$, with **optimal threshold** k^* given by

$$k^* = \ln \frac{P(Y = 0)}{P(Y = 1)}$$

- In some cases, prior probabilities are unknown or cannot be estimated reliably.
- For a classifier ψ , the class specific error rates, $\varepsilon^0[\psi]$ and $\varepsilon^1[\psi]$ are independent of prior probabilities. However, overall error rate $\varepsilon[\psi]$ is not available.
- Also available is optimal threshold D^* .
- Optimal threshold k^* however, is not available.

Minimax classification

- Classifiers have to be defined as a function of k :


$$\psi_k^*(\mathbf{x}) = \begin{cases} 1, & D^*(\mathbf{x}) > k \\ 0, & \text{otherwise} \end{cases} \text{ for } k \in R$$

- Class specific error rates $\varepsilon^0(k) = \varepsilon^0[\psi_k^*]$ and $\varepsilon^1(k) = \varepsilon^1[\psi_k^*]$.
- Each classifier ψ_k^* is a Bayes classifier for a particular set of prior probabilities.
- The Bayes error is a linear combination of $\varepsilon^0(k)$ and $\varepsilon^1(k)$
- Increasing k decreases $\varepsilon^0(k)$ and increases $\varepsilon^1(k)$ and vice versa:
 - A higher k raises the threshold for classifying data as class 1, reducing false positives (fewer class 0 instances misclassified as 1).
 - A higher k makes it harder to classify data as class 1, increasing false negatives (more class 1 instances misclassified as 0).
- A better discriminant should have uniformly smaller $\varepsilon^0(k)$ and $\varepsilon^1(k)$.

Minimax classification

- A plot of $1 - \varepsilon^1(k)$ vs. $\varepsilon^0(k)$ is called **receiver operating characteristics (ROC)** can be used to assess discriminatory power.
- The area under the curve called **area under ROC curve (AUC)** can help evaluate the performance of the discriminant.
- Since ε^* is a linear combination of $\varepsilon^0(k)$ and $\varepsilon^1(k)$, the maximum value the Bayes error can take is $\max\{\varepsilon^0(k), \varepsilon^1(k)\}$. The **minimax criterion** k_{mm} , such that:
$$k_{mm} = \arg \min_k [\max \varepsilon^0(k), \varepsilon^1(k)]$$
- Thus the **minimax classifier**, becomes:

$$\psi_{k_{mm}}^*(\mathbf{x}) = \begin{cases} 1, & D^*(\mathbf{x}) > k_{mm} \\ 0, & \text{otherwise} \end{cases}$$




Bayes Decision Theory

- A general procedure for decision making, where optimal classification is a special case.
- Suppose we are building a classifier for an online medical diagnosis site like WebMD.
- There are three **states of nature** (classes, $c = 3$) a user can be in $Y \in \{y_0, y_1, y_2\}$:
 - 0: Common cold
 - 1: Pneumonia
 - 2: Lung Cancer
- The system can take one of three **actions** $\alpha(x) \in \{\alpha_0, \alpha_1, \alpha_2\} (a = 3)$:
 - Diagnosing the user to have diseases 0, 1, or 2 based on $\alpha_0, \alpha_1, \alpha_2$
 - There could be a case where, $a > c$.
- The user can upload their body temperature x_1 , a pain index x_2 .
- These features are used to compute posterior probabilities $P(Y = j | X = x)$

Bayes Decision Theory

- Now we need to evaluate how well the system works: a **loss function** λ_{ij} which is the cost of taking action α_i (diagnosing disease y_i) when actually the user had a disease y_j .
- We decide on a **loss matrix** as shown in the figure:
 - Diagnosing lung cancer as common cold incurs high penalty.
 - Diagnosing common cold as lung cancer incurs moderate penalty.
- Suppose we have some data from the past based on the reported body temperature what kind of disease patients have had.



Bayes Decision Theory

- We calculate the posterior probabilities using **Bayes Theorem**:

$$P(Y = y_i | X = x) = \frac{P(X = x | Y = y_i)P(Y = Y_i)}{P(X = x)}$$

- Suppose we know these posterior probabilities for a feature x_1 :

$$P(Y = 0 | x_1) = 0.2, P(Y = 1 | x_1) = 0.3, \text{ and } P(Y = 2 | x_1) = 0.5$$

- First, we need to know what would be the risk involved with taking each action, called **conditional risk** $\alpha(x)$:
 - α_0 : $R(\alpha_0) = (0 \cdot 0.2) + (5 \cdot 0.3) + (10 \cdot 0.5) = 6.5$
 - α_1 : $R(\alpha_1) = (2 \cdot 0.2) + (0 \cdot 0.3) + (8 \cdot 0.5) = 4.4$
 - α_2 : $R(\alpha_2) = (4 \cdot 0.2) + (3 \cdot 0.3) + (0 \cdot 0.5) = 1.7$
- Bayes decision rule is that it is enough to select an action with minimum conditional risk, hence $\alpha(x) = \alpha_2$.

Bayes Decision Theory

Generalizing:

- On observing the feature vector $X = x$, and taking an **action** $\alpha(x) \in \{\alpha_0, \alpha_1, \dots, \alpha_{a-1}\}$
- Assuming there are c **states of nature** such that $Y \in \{y_0, y_1, \dots, y_{c-1}\}$.
- And each action incurs a **loss** λ_{ij} = cost of taking action α_i when true state of nature is j .
- Where the action α_i might be simply deciding the true state of nature is i ($a = c$) or $a > c$, where there maybe an extra action of rejecting to take an action.
- The conditional loss upon observing a feature vector $X = x$, called **conditional risk** is given by:

$$R[\alpha(x) = \alpha_i] = \sum_{j=0}^{c-1} \lambda_{ij} P(Y = j \mid X = x)$$

- The **risk** over the entire feature space $x \in R^d$ is given by:

$$R = E[R(\alpha(X))] = \int_{x \in R^d} R(\alpha(x)) p(x) dx$$

- To minimize R , it is enough to select an action $\alpha(x) = \alpha_i$ such that the conditional risk $R[\alpha(x) = \alpha_i]$ is minimum at each value $x \in R^d$. This strategy is called the **Bayes decision rule**, with corresponding optimal **Bayes risk** R^* .

Bayes Decision Theory – Special Case

- Now, if $a = c = 2$:

$$R[\alpha(\mathbf{x}) = \alpha_0] = \lambda_{00}P(Y = 0 | \mathbf{X} = \mathbf{x}) + \lambda_{01}P(Y = 1 | \mathbf{X} = \mathbf{x})$$

$$R[\alpha(\mathbf{x}) = \alpha_1] = \lambda_{10}P(Y = 0 | \mathbf{X} = \mathbf{x}) + \lambda_{11}P(Y = 1 | \mathbf{X} = \mathbf{x})$$

- If $R[\alpha(\mathbf{x}) = \alpha_0] < R[\alpha(\mathbf{x}) = \alpha_1]$, we decide action α_0 if:

$$R[\alpha(\mathbf{x}) = \alpha_1] > R[\alpha(\mathbf{x}) = \alpha_0]$$

$$\lambda_{10}P(Y = 0 | \mathbf{X} = \mathbf{x}) + \lambda_{11}P(Y = 1 | \mathbf{X} = \mathbf{x}) > \lambda_{00}P(Y = 0 | \mathbf{X} = \mathbf{x}) + \lambda_{01}P(Y = 1 | \mathbf{X} = \mathbf{x})$$

$$(\lambda_{10} - \lambda_{00})P(Y = 0 | \mathbf{X} = \mathbf{x}) > (\lambda_{01} - \lambda_{11})P(Y = 1 | \mathbf{X} = \mathbf{x})$$

Applying Bayes Theorem and assuming $\lambda_{10} > \lambda_{00}$ and $\lambda_{01} > \lambda_{11}$:

$$\frac{p(\mathbf{x} | Y = 0)}{p(\mathbf{x} | Y = 1)} > \frac{\lambda_{01} - \lambda_{11}}{\lambda_{10} - \lambda_{00}} \frac{P(Y = 1)}{P(Y = 0)}$$

Bayes Decision Theory – Special Case

- The loss $\lambda_{00} = \lambda_{11} = 0$ and $\lambda_{10} = \lambda_{01} = 1$ is called the **0-1 loss**, equivalent to the misclassification loss.
- With these losses, $\frac{p(x|Y=0)}{p(x|Y=1)} > \frac{\lambda_{01}-\lambda_{11}}{\lambda_{10}-\lambda_{00}} \frac{P(Y=1)}{P(Y=0)}$ becomes $\frac{p(x|Y=0)}{p(x|Y=1)} > \frac{P(Y=1)}{P(Y=0)}$, which is equivalent to the Bayes classifier.
- Thus, if the action is simply deciding the state of nature is $i \forall i \in \{0,1\}$, the 0-1 loss case leads to the **optimal binary classification**.
- The (conditional) risk reduces to (conditional) classification error.
- The Bayes decision rule reduces to Bayes classifier.
- The Bayes risk reduces to Bayes error.

Thank You